

---

---

# CS 301

# High-Performance Computing

---

---

## Lab Report-02

---

## Matrix Multiplication

---

## Performance Analysis

---

Tanish Patel (202301411)  
Neelabh Rana (202301476)

# Contents

<b>1</b>	<b>Introduction</b>	<b>3</b>
<b>2</b>	<b>Hardware Details</b>	<b>3</b>
2.1	LAB207 PCs . . . . .	3
2.2	HPC Cluster . . . . .	3
<b>3</b>	<b>Methodology</b>	<b>4</b>
<b>4</b>	<b>Matrix Multiplication</b>	<b>4</b>
4.1	Problem Size vs Time . . . . .	4
4.1.1	Local System . . . . .	4
4.1.2	HPC Cluster . . . . .	4
<b>5</b>	<b>Transpose Method</b>	<b>7</b>
<b>6</b>	<b>Block Matrix Multiplication</b>	<b>7</b>
<b>7</b>	<b>Observations</b>	<b>7</b>
<b>8</b>	<b>Conclusion</b>	<b>8</b>

# 1 Introduction

This report analyzes the performance of matrix multiplication algorithms on local and HPC cluster systems. We compare different loop orderings, transpose-based optimization, and block-based multiplication.

Performance is evaluated using execution time for problem sizes ranging from  $2^1$  to  $2^{12}$ .

## 2 Hardware Details

### 2.1 LAB207 PCs

- Architecture: x86\_64
- CPU(s): 12
- Threads per core: 2
- Cores per socket: 6
- Processor: Intel Core i5-12500
- L1 Cache: 288K
- L2 Cache: 7.5M
- L3 Cache: 18M

### 2.2 HPC Cluster

- Architecture: x86\_64
- CPU(s): 16
- Threads per core: 1
- Cores per socket: 8
- Sockets: 2
- Processor: Intel Xeon E5-2640 v3
- L1 Cache: 32K
- L2 Cache: 256K
- L3 Cache: 20480K

### 3 Methodology

Matrix multiplication is implemented using:

- Six loop orderings
- Transpose optimization
- Block multiplication

Execution time is measured and converted to logarithmic scale:

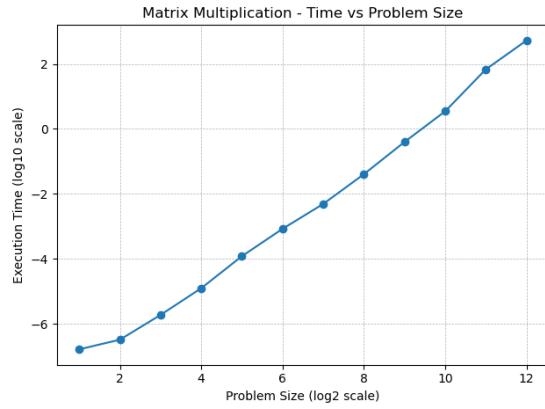
$$C[i][j] = \sum_{k=0}^{n-1} A[i][k] \times B[k][j] \quad (1)$$

## 4 Matrix Multiplication

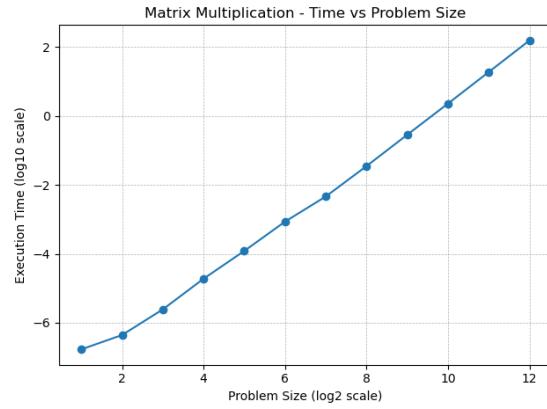
### 4.1 Problem Size vs Time

#### 4.1.1 Local System

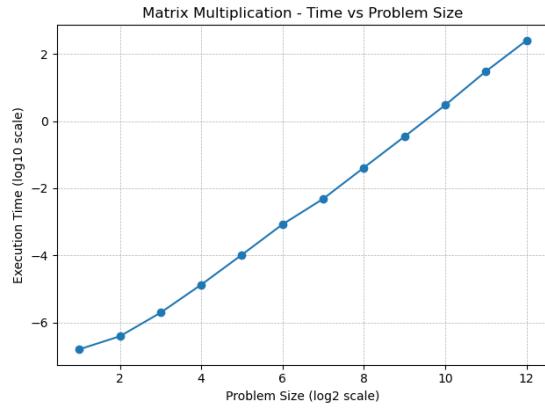
#### 4.1.2 HPC Cluster



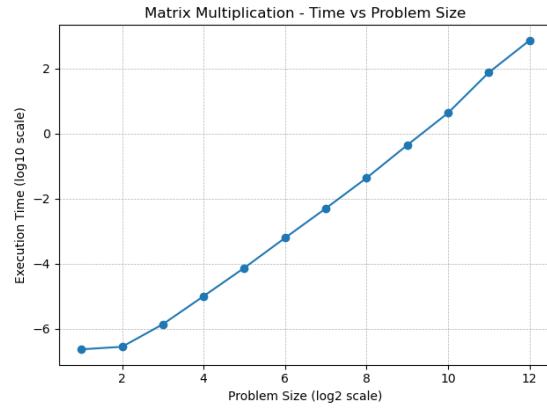
(a) IJK



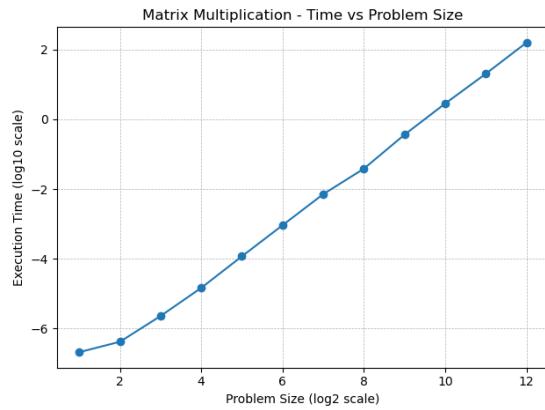
(b) IKJ



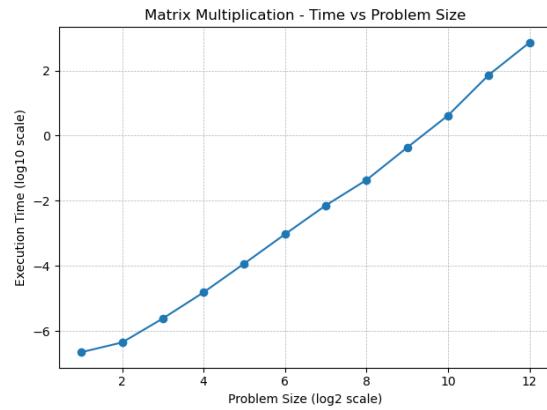
(c) JIK



(d) JKI



(e) KIJ



(f) KJI

Figure 1: Execution Time vs Problem Size (Local System)

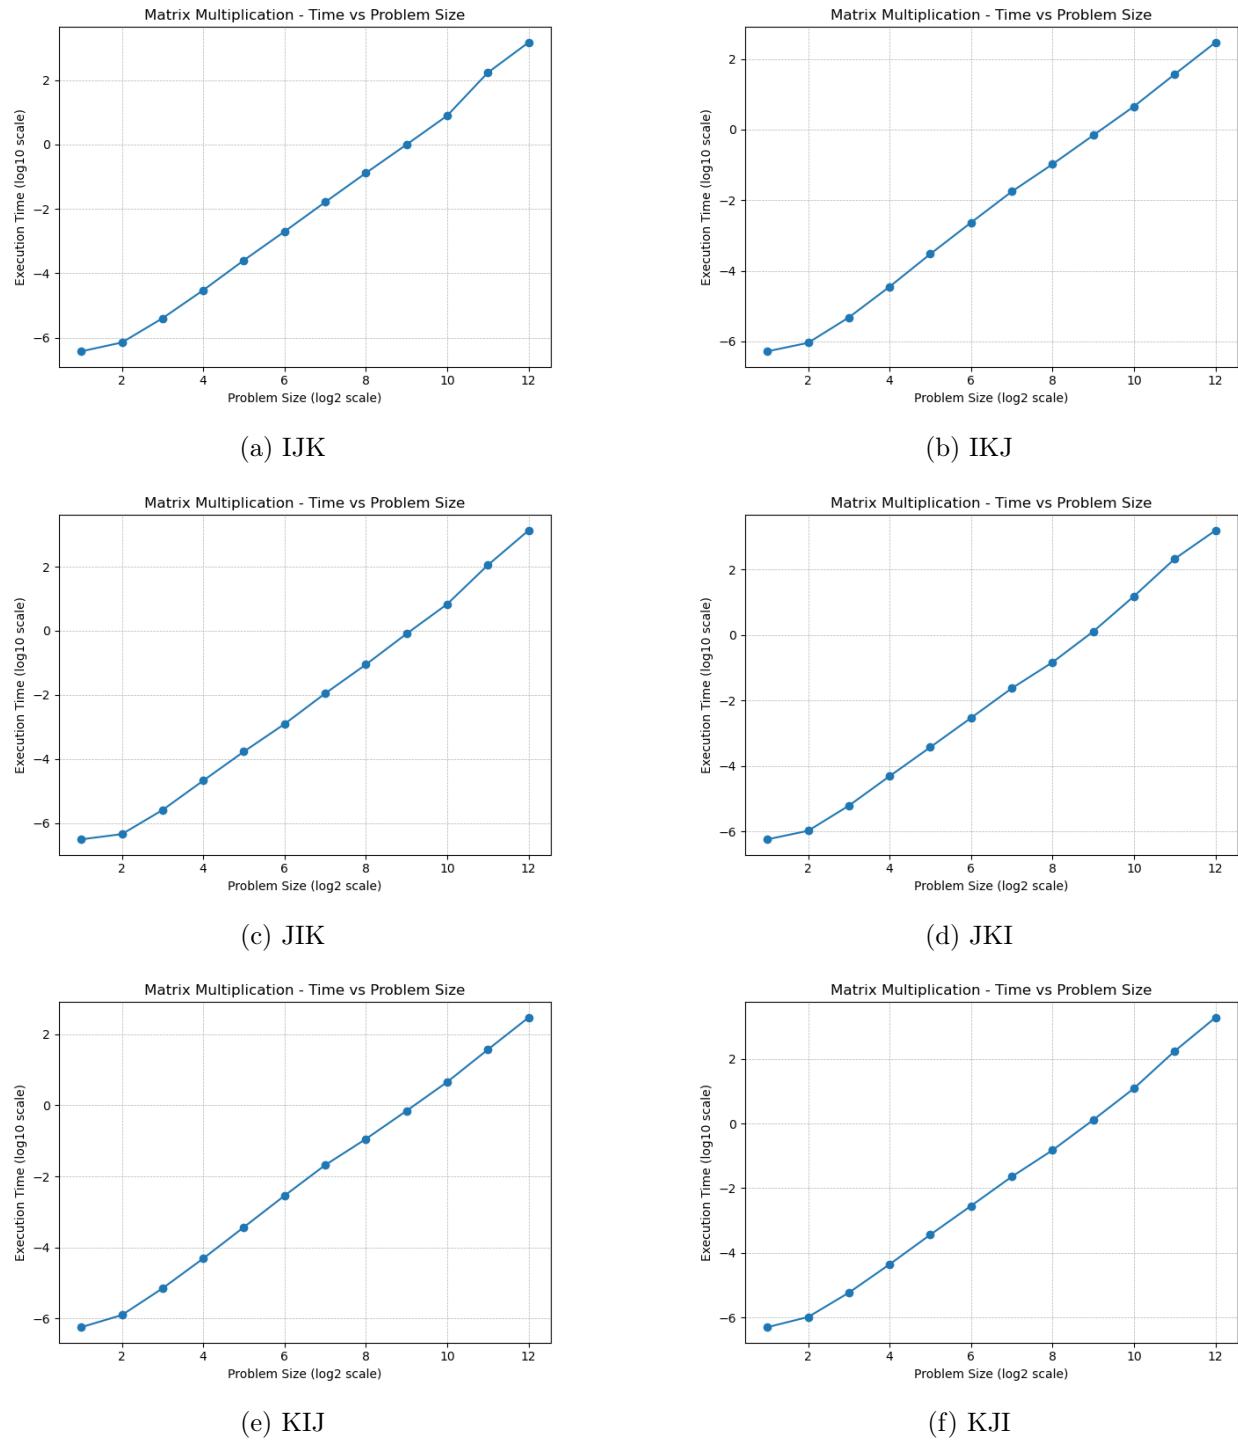


Figure 2: Execution Time vs Problem Size (HPC Cluster)

## 5 Transpose Method

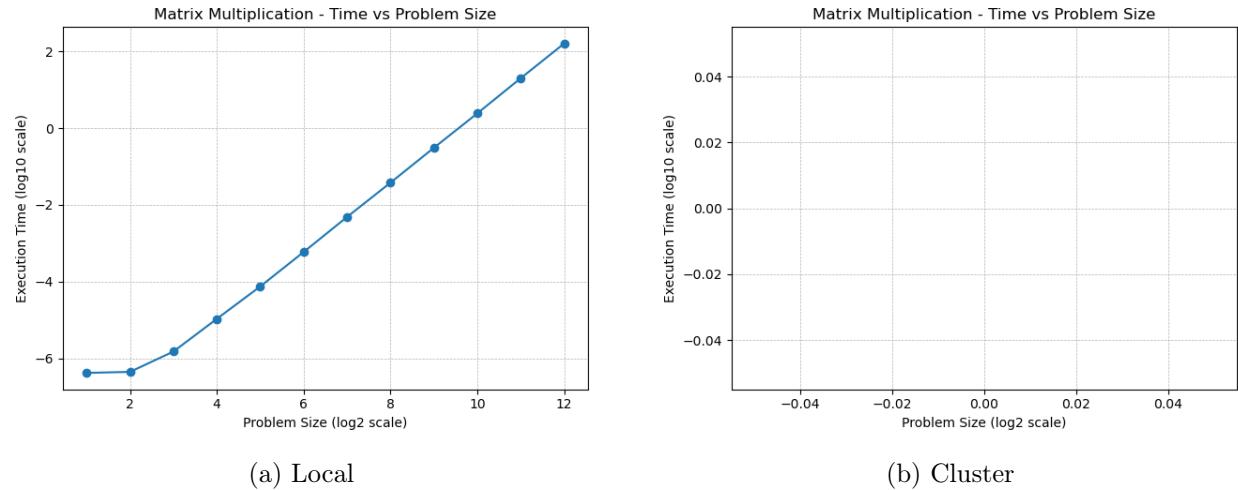


Figure 3: Transpose Method Performance

## 6 Block Matrix Multiplication

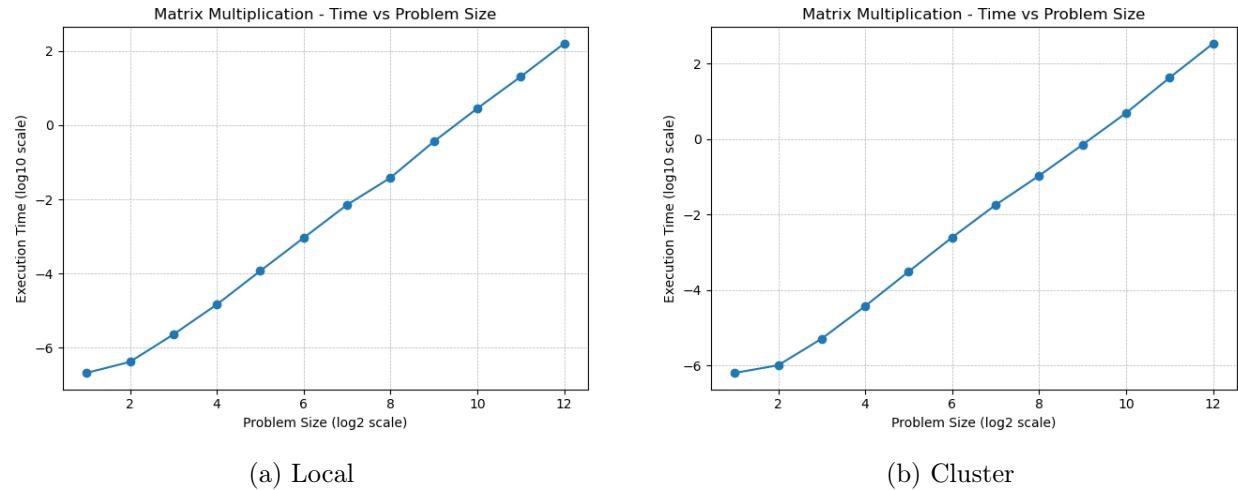


Figure 4: Block Method Performance

## 7 Observations

- Loop ordering strongly affects cache performance.
- KIJ and IKJ provide best locality.

- Transpose improves memory access.
- Blocking enhances cache reuse.

## 8 Conclusion

Memory access patterns play a crucial role in matrix multiplication performance. Loop optimization, transposition, and blocking significantly improve runtime, especially on HPC systems.